

Design Specifications for “Hosting of Virtual-labs using the one-lab-per-vm model”

Suraj Samal

05 November 2013

Contents

1	Introduction	1
2	Document Revision	1
3	Basic Architecture	2
3.1	Overview	2
3.2	Actors	2
3.2.1	Lab Developer	2
3.2.2	Lab Administrator	2
3.2.3	Lab User	2
3.3	Entities	2
3.3.1	LabDepository	2
3.3.2	Lab	4
3.3.3	LabManager	4
3.3.4	Host	5
3.3.5	VMManager	5
3.3.6	VM	6
3.4	Relationships	7
3.4.1	LabDepository - repository - revision	7
3.4.2	Lab - repository - revision	7
3.4.3	LabManager - host - vmmgr - vm - lab	7
3.5	Workflows	7
3.5.1	Lab Developer Workflows	7
3.5.2	Lab Administrator Workflows	8
3.5.3	User Workflows	9
3.5.4	View a Lab	9
3.5.5	Other Implicit Workflows	9
3.5.6	Log Lab Information	9
3.5.7	AutoPurge Lab History	9

4	Components and Interfaces	9
4.1	Lab Manager	9
4.2	VM Manager	10
4.3	DeveloperPortal	11
4.4	DeploymentDashboard	11
4.5	LabInfoDatabase	11
5	Network Architecture	11
6	Security Architecture	11
7	Performance Model	12
8	Reliability and Availability Model	12
9	Backup Model	12
10	Scalability Model	12

List of Tables

List of Figures

1 Introduction

The document discusses the design of the overall architecture of the hosting of virtual-labs using the one-lab-per-vm model. This is as per the requirements specified in “Minutes of the 2013-07-25 Thu Expert Committee meeting evaluating VLEAD’s progress in virtual lab integration” document at **Section-4, Item-3**

VLEAD (Virtual Labs Engineering and Architecture Division) team was setup in June 2012 as a central engineering team for integrating all the virtual-labs (around 180 in number) across all disciplines and institutes onto a common data-center (currently located at IIIT Hyderabad). Currently(as of 2013-11-01) around 100 labs are version-controlled and around 50 hosted out of IIIT data-centre.

2 Document Revision

Current Revision 0.1

Revision Date 2013-11-04

3 Basic Architecture

3.1 Overview

Below is an overview of the overall system describing all the actors, entities and their interfaces: [file:overview.jpg](#)]]

3.2 Actors

3.2.1 Lab Developer

An person who has agreed to use the services of VLEAD as per the **terms of association** and follows certain standard processes to maintain his/her lab during its development life-cycle. In specific, the roles are as follows:

- Checkin the lab contents (sources,dependencies, scripts and other files) into a lab-depository.
- Keep updating the lab-depository with newer revisons of lab contents.
- Instantiate a test lab-instance for testing and debugging issues.
- Instantiate a live lab-instance.
- View live lab-instance statistics.

3.2.2 Lab Administrator

An actor who is responsible for administering all the hosted labs. In specific, the roles are as follows:

- Allocate a unique labid and a depository(collection of repositories) to a lab
- Allocation of resources(physical machines,ip address pools, vmid pools) to the lab-manager and vmmanager

3.2.3 Lab User

These are end-users who use the virtual-labs and its experiments

3.3 Entities

3.3.1 LabDepository

All labs are allocated a unique-id and a lab-depository by the labs administrator. A lab-depository represents a collection of various repositories associated with a lab.

lab-depository - An **Object** describing the property of all repositories of a particular lab

labid - Unique identifier of the lab

labinfo - **Object** describing basic properties of a lab

labinst - One of the defined **enumerations** (IITB, IITK, IIITH ,,,)

labdisc - One of the defined **enumerations** (chemical, mechanical)

labos - **Object** describing a particular operating-system version

osname - Name of the operating system

osversion - Specific version of the operating system

repos - Collection of repositories

metadata - A structured **object** representation of depository contents describing the number of repos present, actual repos present, their type . This repository is regenerated everytime the lab-developer makes a commit to other repositories.

numrepos - Total number of repositories present

reposit1 - Identification of each repositories

reposit1 - **reposit2** -

reposit1 - .

reposit1 - .

reposit1 - **repositN** -

repo1 - A repository **object** which refers to a svn, git or bsr repository

reposit - Identification text that can be used to checkout the repository.
(Eg: cse01, mech09)

reponame - Display text (Eg: Frontend, Backend, UI etc)

repositype - One of the supported **enumerated** types - (git, svn, bsr)

revsnum - Number of revisions of the repository (Eg: 20)

rev - **Object** defining a particular repository revision

revno - Unique revision number generated by the repository tool. (Eg: 10)

date - Date/Time the revision was checked into the repository. (Eg: 2013-11-10 16:30)

user - Text representing user who checked the revision. (Eg: ramakrishna)

diskspace - Approximate disk-space required. (Eg: 30G)

ram - Approximate memory required. (Eg: 256M)

staticdeps - An **object** describing a list of packages the lab depends on. (Eg: apache2, opencv)

dep1

dep2 . .

depn

runtimedeps - An **object** describing a list of services to be enabled/started. Services may mean standard packages (eg. apache2) or other custom made scripts (Eg: backup) to be configured during installation of the lab.

dep1

dep2 . . .

depn

size - Number representing the size of the particular repository revision (**Optional**)

repo2 -

repo2 - .

repo2 - .

repoN -

3.3.2 Lab

An instance of a lab (inactive) which refers to a complete set of properties that can be used to instantiate a particular lab revision. All these properties can be loaded directly from the lab-depository by using its unique labid, unique repoid and a unique revision no.

lab - **Object** describing an lab

labid - Unique id to identify the lab from others

labinfo - **Object** describing basic properties of a lab

repo - **Object** describing a particular repository of a lab

rev - **Object** describing a particular revision of a particular repository of a lab

3.3.3 LabManager

An entity that monitors a set of physical hosts, accepts requests for creation, modification and deletion of lab-instances and sends request to appropriate vm-manager for life-cycle management of labinstances

labmanager - An entity responsible for managing the various vm-managers

labmanagerid - Unique id to describe a labmanager

hosts - **Object** representation of a list of physical-hosts

host1 - **Object** representation of a physical host (described later) . . .

host2 - . . .

host3 -

runtime runtime characteristics of the labmanager

start_{time} - timestamp the labmanager was instantiated

3.3.4 Host

A physical host entity managed by a lab-manager and hosting a single vm-manager

Host - Entity representing a physical host

hostname - Common name of the host

vmmgr - **Object** representation of the vm-manager (described later) managing the host

hostid - Unique-id representation of the host

hostip - IPaddress of the physical host

resource - **Object** representation of resources of the physical host

diskspace - (Eg. 2000GB)

mem - (Eg. 64GB)

cpu - (Eg. 2)

runtime - Runtime properties of the host

status - one of running, stopped, shutoff

starttime - timestamp the host was started

useddiskspace - (Eg. 100GB)

usedmem - (Eg. 20GB)

usedcpu - (Eg. 1)

3.3.5 VMManager

An entity that is responsible for managing virtual machines(vms) on a particular host

vmmgr - Entity describing an instance of a vm-manager residing on a physical machine

vmmgrid - Unique id to represent the vm-manager

vms - List of vm objects

vm1 - **Object** representation of a vm (described later)

vm2 -

vmN -

resources - **Object** representation of resources

vmids - List of available vmids

vmid1 -

vmid2 - .

vmidn -

ips - List of available ips

ip1 -
ip2 - . .
ipn -
runtime - Runtime properties
status - up, down, stopped
start_{time} - start timestamp

3.3.6 VM

A VM is a running instance of a lab.

vm - An active instance of a lab that runs on a specified host

guid - Global Universal id of the vm generated to identify the

vmid - Unique identification of a vm among its current running VMs. This is allocated from a defined pool of ids when the vm is created and re-sent to the pool when the vm gets destructed.

vmname - Common name to identify the VM instance.

vmos - Operating system **object** of the running vm.

osname - Name of the operating system

osversion - Particular version of the operating system

lab - A particular instance of a lab associated with a vm

runtime - **Object** describing run-time properties of the vm

state - running, stopped, suspended, archived

createddate - Creation time-stamp of the VM

modifieddate - Modification time-stamp of the VM

lastbackedup - Timestamp when the vm was last backedup

stats - **Object** describing stats of a vm

userstats - User-level statistics of the vm

userinfo -

perfstats - **cpuinfo** -

meminfo -

netinfo -

3.4 Relationships

3.4.1 LabDepository - repository - revision

[Lab-Depository] 1 ————— * [repo] 1 ————— * [rev]

3.4.2 Lab - repository - revision

[Lab] 1 ——— 1 [repo] 1 ——— 1 [rev]

3.4.3 LabManager - host - vmmgr - vm - lab

[Labmanager] * ——— * [host] 1 ——— 1 [vmmgr] 1 ——— * [vm] 1 ——— 1 [lab]

3.5 Workflows

3.5.1 Lab Developer Workflows

- Create a Lab

[./Create-a-Lab.jpg](#)

- Update a Lab

[./Update-a-lab.jpg](#)

- Test a Lab

[./Test-a-Lab.jpg](#)

- Release a Lab

[./Release-a-Lab.jpg](#)

- Delete a Lab



- Fetch Lab-Statistics

[./Fetch-lab-statistics.jpg](#)

3.5.2 Lab Administrator Workflows

- Create a Lab Repository
- Delete a Lab Repository
- Update Resource Information
 - Physical Machine Resources
 - Network Parameters
 - VM Manager Information
- Update Lab Backup Schedule
- Take a Lab run-time snapshot
- Restore a Lab from its snapshot backup
- Deactivate a Lab
- Monitor VM Statistics
- Modify VM Run-time Parameters
- Purge a VM
- Purge VM logs

3.5.3 User Workflows

3.5.4 View a Lab

3.5.5 Other Implicit Workflows

3.5.6 Log Lab Information

3.5.7 AutoPurge Lab History

4 Components and Interfaces

- Following are the components that need to be designed for the proposed architecture:

4.1 Lab Manager

LabOperator Manages basic operations for the life-cycle management of lab

- createLab(vmmanager, lab)
- updateLab(vmmanager, lab)
- deleteLab(vmmanager, lab)
- updateresources() - Adds or removes resources information (Eg. vmmanager, hosts)

LabMonitor Regularly monitors the status of labs and vms

- ping(vmmanager, lab)

LabLogger Logs status and history information to the lab-info database

- loginfo()
- logwarn()
- logerror()
- purgelogs()

LabStatsCollector • collectvmstats(vmmanager)

- collectlabstats(vmmanager, lab)
- collectrepostats(lab)
- updatevmstatstoDB()
- updatelabstatstoDB()

BackupManager • backup(vmmanager, lab)

- restore(vmmanager, lab)
- schedule(lab)

4.2 VM Manager

VMOperator Manages basic operations for life-cycle of a vm and a lab

- createvm(lab)
- updatevm(vmid)
- deletevm(vmid)
- stopvm(vmid)
- startvm(vmid)
- updateresources(host)
- checkoutlab(vmid, lab)
- buildlab(vmid, lab)
- deploylab(vmid, lab)
- activatelab(vmid, lab)
- testlab(vmid, lab)
- restorelab(lab, snapshot)
- backuplab (lab, snapshot)
- updateinfoDB()

VMMonitor • pinglab(lab)

- getcpuinfo(vmid)
- getmemusage(vmid)
- getnetworkusage(vmid)
- getuserstats(vmid)
- getcpuinfo(host)
- getmemusage(host)
- getnetworkusage(host)

VMLogger • loginfo()

- logwarn()
- logerror()
- purgelogs()

CommandsGenerator A component that generates the configuration commands based on operation specified by the VMOperator

- generateconfig(configid)

CommandExecutor A component that runs the configuration commands generated earlier by the CommandsGenerator

- applyconfig(configid)

4.3 DeveloperPortal

- createdepository(lab)
- createrepository(labdepository, lab)
- updaterepository(labdepository, lab)
- deleterepostitory(labdepository, l
- deletedepository(lab)
- sendrequest(labmanager, lab, operation) - Operation could be one of create/update/test/release a lab or getlabstats
- updateresources(labmanager) - Information about physical-hosts, network parameters etc

4.4 DeploymentDashboard

- getlabsStatus(labmanager)
- getlabsHistory(labInfoDb)

4.5 LabInfoDatabase

- VMHistory
- LabHistory
- VMManagerHistory
- LabManagerHistory

5 Network Architecture

Presented below is a network architecture diagram of the proposed solution: [Network](#)

6 Security Architecture

- Firewall rules are configured at the router-interface for translating public requests to private requests.
- Labs are accessed by users through a web-proxy that logically isolates the actual lab-instances from public world. In any case, the security of the web-proxy host is compromised. The web-proxy can be configured for additional security and monitored for user statistics. Additionally, only specific ports are enabled so that the labs can be accessed over web.

- Labs are accessed by lab-developers using a gateway that isolates the actual lab-vms from the public world. Additionally, the lab-vms are proposed to be in a separate sub-network for additional security.

7 Performance Model

8 Reliability and Availability Model

9 Backup Model

- All labs would be backup

10 Scalability Model