

Design Specifications for “Hosting of Virtual-labs using the one-lab-per-vm model”

Suraj Samal

05 November 2013

Contents

1	Introduction	1
2	Document Revision	1
3	Basic Architecture	2
3.1	Overview	2
3.2	Actors	2
3.2.1	Lab Developer	2
3.2.2	Lab Administrator	2
3.2.3	Lab User	2
3.3	Entities	2
3.3.1	LabDepository	2
3.3.2	Lab	4
3.3.3	LabManager	4
3.3.4	Host	4
3.3.5	VMManager	5
3.3.6	VM	6
3.4	Relationships	6
3.4.1	LabDepository - repository - revision	6
3.4.2	Lab - repository - revision	6
3.4.3	LabManager - host - vmmgr - vm - lab	7
3.5	Workflows	7
3.5.1	Lab Developer Workflows	7
3.6	Lab Administrator Workflows	8
3.7	User Workflows	8
3.7.1	View a Lab	8
3.8	Other Implicit Workflows	8
3.8.1	Log Lab Information	8
3.8.2	AutoPurge Lab History	8

4	Components and Interfaces	8
4.1	Lab Manager	8
4.2	VM Manager	8
4.3	DeveloperPortal	9
4.4	DeploymentDashboard	9
4.5	LabInfoDatabase	9
4.6	9
5	Network Architecture	9
6	Security Architecture	9
7	Performance Model	9
8	Reliability and Availability Model	9
9	Backup Model	9
10	Scalability Model	9

List of Tables

List of Figures

1 Introduction

The document discusses the design of the overall architecture of the hosting of virtual-labs using the one-lab-per-vm model. This is as per the requirements specified in “Minutes of the 2013-07-25 Thu Expert Committee meeting evaluating VLEAD’s progress in virtual lab integration” document at **Section-4, Item-3**

VLEAD (Virtual Labs Engineering and Architecture Division) team was setup in June 2012 as a central engineering team for integrating all the virtual-labs (around 180 in number) across all disciplines and institutes onto a common data-center (currently located at IIT Hyderabad). Currently(as of 2013-11-01) around 100 labs are version-controlled and around 50 hosted out of IIT data-centre.

2 Document Revision

Current Revision 0.1

Revision Date 2013-11-04

3 Basic Architecture

3.1 Overview

Below is an overview of the overall system describing all the actors, entities and their interfaces:

Architecture

3.2 Actors

3.2.1 Lab Developer

An person who has agreed to use the services of VLEAD as per the **terms of association** and follows certain standard processes to maintain his/her lab during its development life-cycle. In specific, the roles are as follows:

- Checkin the lab contents (sources,dependencies, scripts and other files) into a lab-depository.
- Keep updating the lab-depository with newer revisons of lab contents.
- Instantiate a test lab-instance for testing and debugging issues.
- Instantiate a live lab-instance.
- View live lab-instance statistics.

3.2.2 Lab Administrator

An actor who is responsible for administering all the hosted labs. In specific, the roles are as follows:

- Allocate a unique labid and a depository(collection of repositories) to a lab
- Allocation of resources(physical machines,ip address pools, vmid pools) to the lab-manager and vmmanager

3.2.3 Lab User

These are end-users who use the virtual-labs and its experiments

3.3 Entities

3.3.1 LabDepository

All labs are allocated a unique-id and a lab-depository by the labs administrator. A lab-depository represents a collection of various repositories associated with a lab.

lab-depository - An **Object** describing the property of all repositories of a particular lab

labid - Unique identifier of the lab

labinfo - **Object** describing basic properties of a lab

labinst - One of the defined **enumerations** (IITB, IITK, IIITH ,...)

labdisc - One of the defined **enumerations** (chemical, mechanical ,...)

labos - **Object** describing a particular operating-system version

osname - Name of the operating system

osversion - Specific version of the operating system

labmetadata - A structured **object** representation of depository contents describing the number of repos present, actual repos present, their type . This is regenerated everytime the lab-developer makes a commit to a repository.

numrepos - Sum of all repositories present in the repository

repo - A repository **object** which refers to a svn, git or bzt repository

repositoid - Identification text that can be used to checkout the repository.
(Eg: cse01, mech09)

reponame - Display text (Eg: Frontend, Backend, UI etc)

repotype - One of the supported **enumerated** types - (git, svn, bzt)

revsnum - Number of revisions of the repository (Eg: 20)

rev - **Object** defining a particular repository revision

revno - Unique revision number generated by the repository tool. (Eg: 10)

date - Date/Time the revision was checked into the repository. (Eg: 2013-11-10 16:30)

user - Text representing user who checked the revision. (Eg: ramakrishna)

diskspace - Approximate disk-space required. (Eg: 30G)

ram - Approximate memory required. (Eg: 256M)

staticdeps - An **object** describing a list of packages the lab depends on. (Eg: apache2, opencv)

dep1

dep2 . .

depn

runtimedeps - An **object** describing a list of services to be enabled/started. Services may mean standard packages (eg. apache2) or other custom made scripts (Eg: backup) to be configured during installation of the lab.

dep1

dep2 . . .

depn

size - Number representing the size of the particular repository revision (**Optional**)

3.3.2 Lab

An instance of a lab (inactive) which refers to a complete set of properties that can be used to instantiate a particular lab revision. All these properties can be loaded directly from the lab-depository by using its unique labid, unique repoid and a unique revision no.

lab - **Object** describing an lab

labid - Unique id to identify the lab from others

labinfo - **Object** describing basic properties of a lab

repo - **Object** describing a particular repository of a lab

rev - **Object** describing a particular revision of a particular repository of a lab

3.3.3 LabManager

An entity that monitors a set of hosts, accepts requests for creation, modification and deletion of labinstances and sends request to appropriate vm-manager for life-cycle management of labinstances

labmanager - An entity responsible for managing the various vm-managers

labmanagerid - Unique id to describe a labmanager

hosts - **Object** representation of a list of physical-hosts

host1 - **Object** representation of a physical host (described later) . . .

host2 - . . .

host3 -

runtime runtime characteristics of the labmanager

start_{time} - timestamp the labmanager was instantiated

3.3.4 Host

A physical host entity managed by a lab-manager and hosting a single vm-manager

Host - Entity representing a physical host

hostname - Common name of the host

vmmgr - **Object** representation of the vm-manager (described later) managing the host

hostid - Unique-id representation of the host

hostip - IPaddress of the physical host

resource - **Object** representation of resources of the physical host

diskspace - (Eg. 2000GB)

mem - (Eg. 64GB)

cpu - (Eg. 2)

runtime - Runtime properties of the host

status - one of running, stopped, shutoff

start_{time} - timestamp the host was started

useddiskspace - (Eg. 100GB)

usedmem - (Eg. 20GB)

usedcpu - (Eg. 1)

3.3.5 VMManager

An entity that is responsible for managing virtual machines(vms) on a particular host

vmmgr - Entity describing an instance of a vm-manager residing on a physical machine

vmmgrid - Unique id to represent the vm-manager

vms - List of vm objects

vm1 - **Object** representation of a vm (described later)

vm2 -

vmn -

resources - **Object** representation of resources

vmids - List of available vmids

vmid1 -

vmid2 - .

vmidn -

ips - List of available ips

ip1 -
ip2 - . .
ipn -
runtime - Runtime properties
status - up, down, stopped
start_{time} - start timestamp

3.3.6 VM

A VM is a running instance of a lab.

vm - An active instance of a lab that runs on a specified host

guid - Global Universal id of the vm generated to identify the

vmid - Unique identification of a vm among its current running VMs. This is allocated from a defined pool of ids when the vm is created and re-sent to the pool when the vm gets destructed.

vmname - Common name to identify the VM instance.

vmos - Operating system **object** of the running vm.

osname - Name of the operating system

osversion - Particular version of the operating system

lab - A particular instance of a lab associated with a vm

runtime - **Object** describing run-time properties of the vm

state - running, stopped, suspended, archived

createddate - Creation time-stamp of the VM

modifieddate - Modification time-stamp of the VM

lastbackedup - Timestamp when the vm was last backedup

stats - **Object** describing stats of a vm

userstats - User-level statistics of the vm

userinfo -

perfstats - **cpuinfo** -

meminfo -

netinfo -

3.4 Relationships

3.4.1 LabDepository - repository - revision

[Lab-Depository] 1 ————— * [repo] 1 ————— * [rev]

3.4.2 Lab - repository - revision

[Lab] 1 ——— 1 [repo] 1 ——— 1 [rev]

3.4.3 LabManager - host - vmmgr - vm - lab

[Labmanager] * ——— * [host] 1 ——— 1 [vmmgr] 1 ——— * [vm] 1 ——— 1 [lab]

3.5 Workflows

3.5.1 Lab Developer Workflows

- Create a Lab
- Update a Lab
- Test a Lab
- Release a Lab
- Delete a Lab
- Fetch Lab-Statistics
- Create a Lab Repository
- Delete a Lab Repository
- Update Resource Information
 - Physical Machine Resources
 - Network Parameters

- VM Manager Information
- Update Lab Backup Schedule
- Take a Lab run-time snapshot
- Restore a Lab from its snapshot backup
- Deactivate a Lab
- Monitor VM Statistics
- Modify VM Run-time Parameters
- Purge a VM
- Purge VM logs

3.6 Lab Administrator Workflows

3.7 User Workflows

3.7.1 View a Lab

3.8 Other Implicit Workflows

3.8.1 Log Lab Information

3.8.2 AutoPurge Lab History

4 Components and Interfaces

- Following are the components that need to be designed for the proposed architecture:

4.1 Lab Manager

- LabOperator
- LabMonitor

- LabLogger
- LabStatsCollector

4.2 VM Manager

- VMOperator
- VMMonitor
- CommandsGenerator
- CommandExecutor
- VMLogger
- VMStatsCollector

4.3 DeveloperPortal

4.4 DeploymentDashboard

4.5 LabInfoDatabase

- VMStats

4.6

5 Network Architecture

Presented below is a network architecture diagram of the proposed solution:

Network

6 Security Architecture

- Firewall rules are configured at the router-interface for translating public requests to private requests.
- Labs are accessed by users through a web-proxy that logically isolates the actual lab-instances from public world. In any case, the security of the web-proxy host is compromised. The web-proxy can be configured for additional security and monitored for user statistics. Additionally, only specific ports are enabled so that the labs can be accessed over web.
- Labs are accessed by lab-developers using a gateway that isolates the actual lab-vms from the public world. Additionally, the lab-vms are proposed to be in a separate sub-network for additional security.

- 7 Performance Model**
- 8 Reliability and Availability Model**
- 9 Backup Model**
- 10 Scalablility Model**